

Image Pyramids

Finding Waldo

- Let's revisit the problem of finding Waldo
- This time he is on the road



image



template (filter)

Finding Waldo

- He comes closer but our filter doesn't know that
- How can we find Waldo?



image



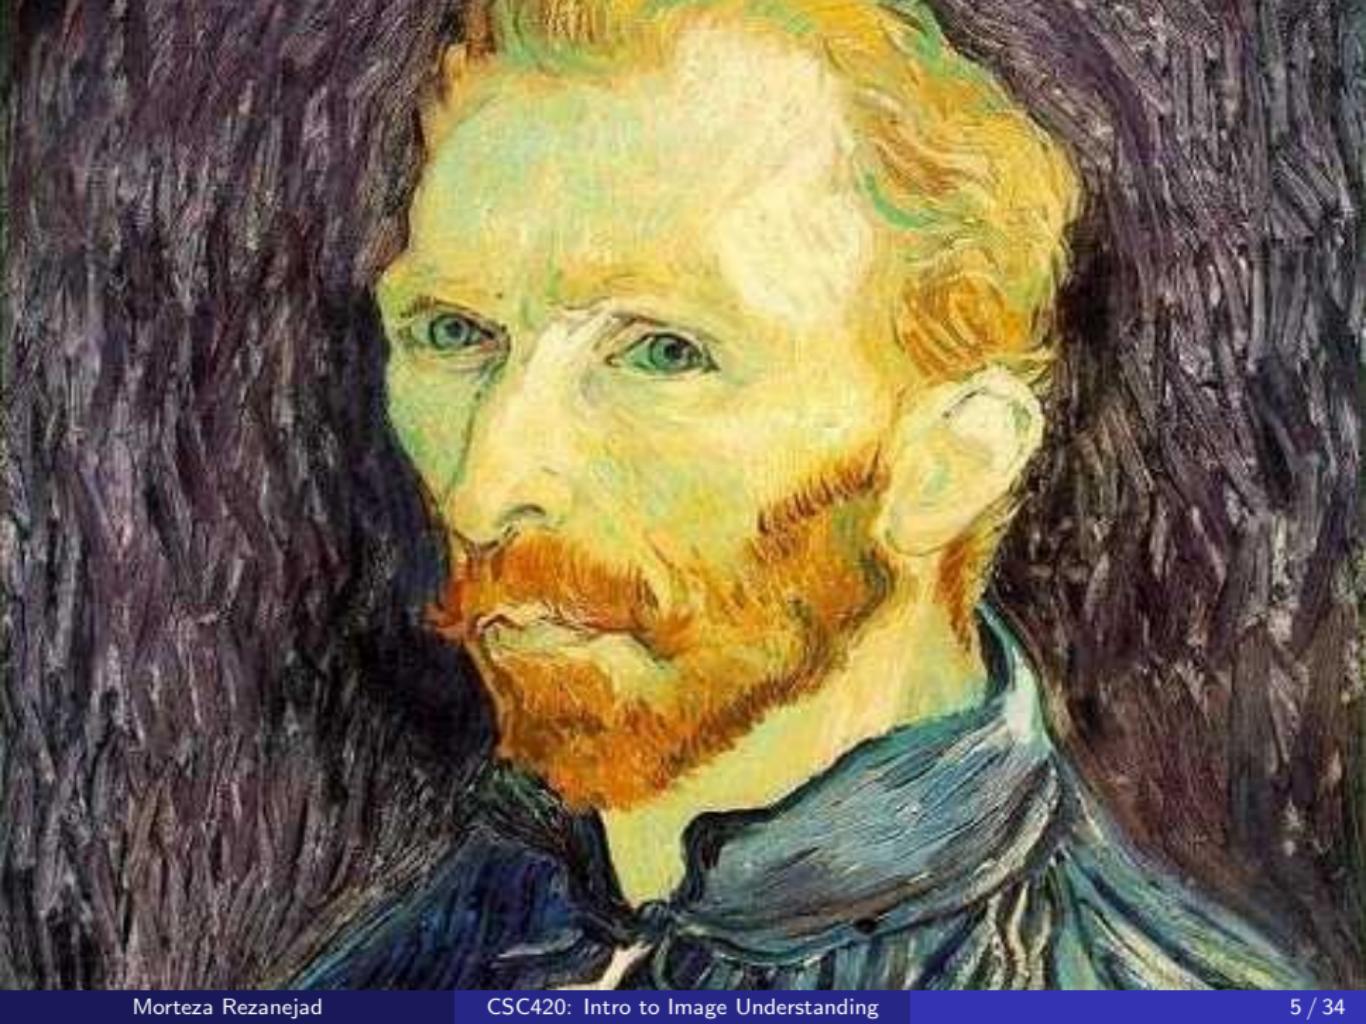
template (filter)

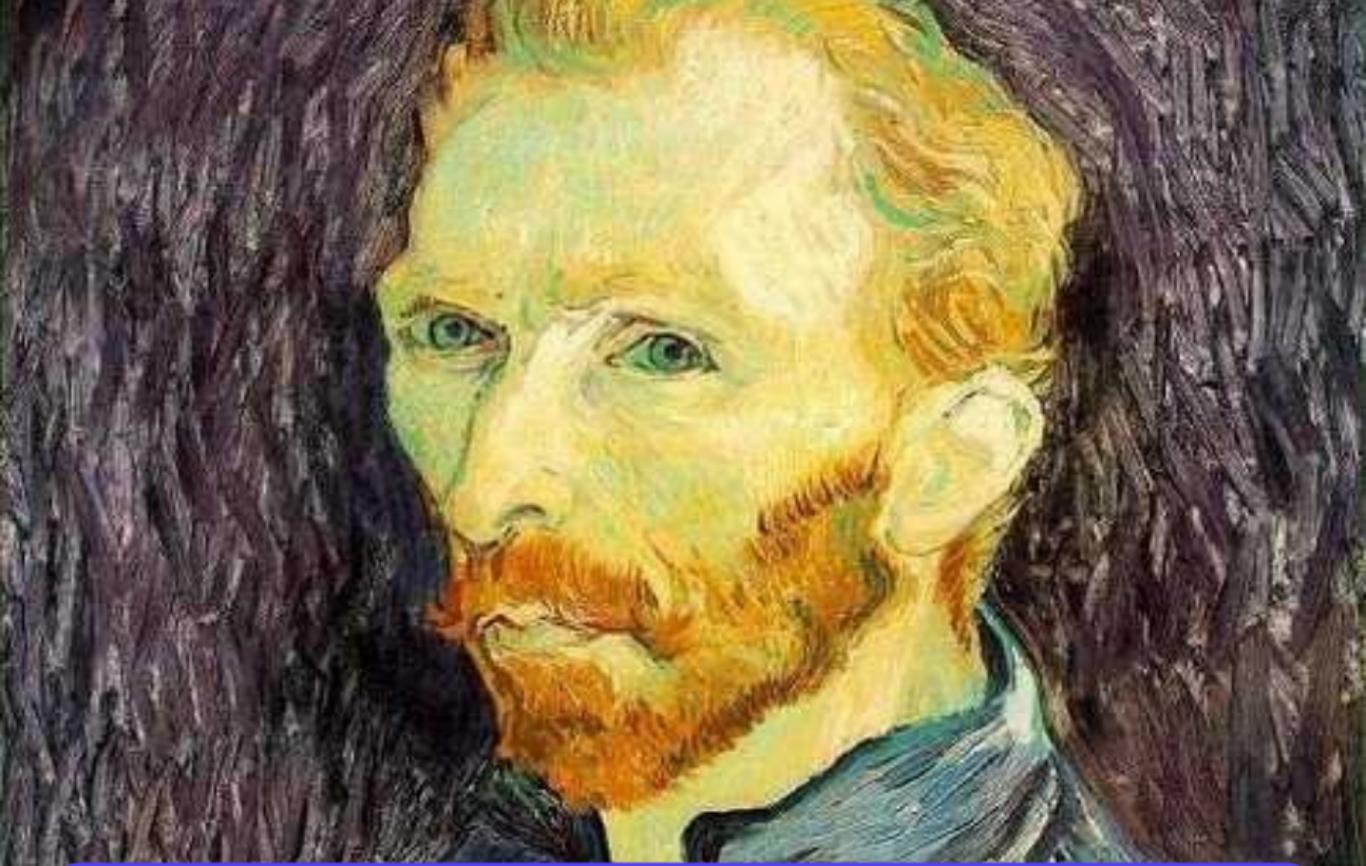
Idea: Re-size Image

- Re-scale the image multiple times! Do correlation on every size!



template (filter)

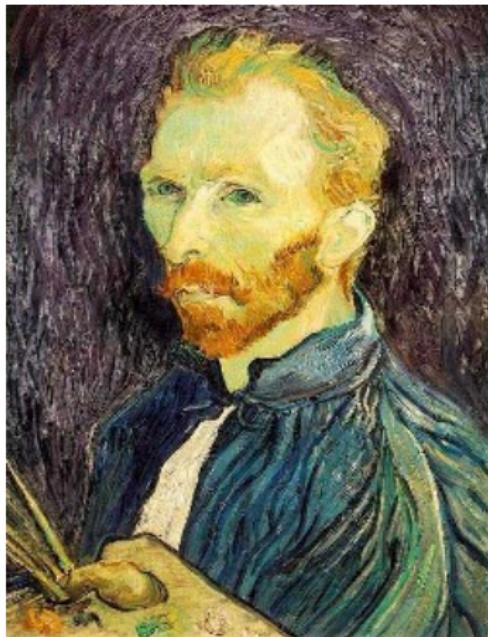




This image is huge. How can we make it smaller?

Image Sub-Sampling

- **Idea:** Throw away every other row and column to create a $1/2$ size image



1/4

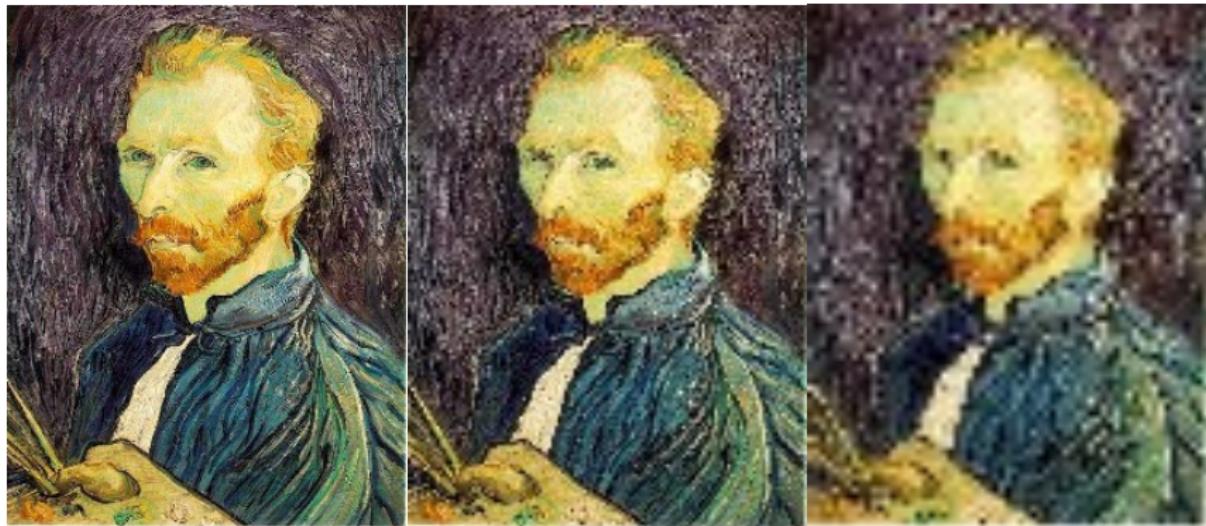


1/8

[Source: S. Seitz]

Image Sub-Sampling

- Why does this look so crulty?



1/2

1/4 (2x zoom)

1/8 (4x zoom)

[Source: S. Seitz]

Even worse for synthetic images

- I want to resize my image by factor 2
- And I take every other column and every other row (1st, 3rd, 5th, etc)

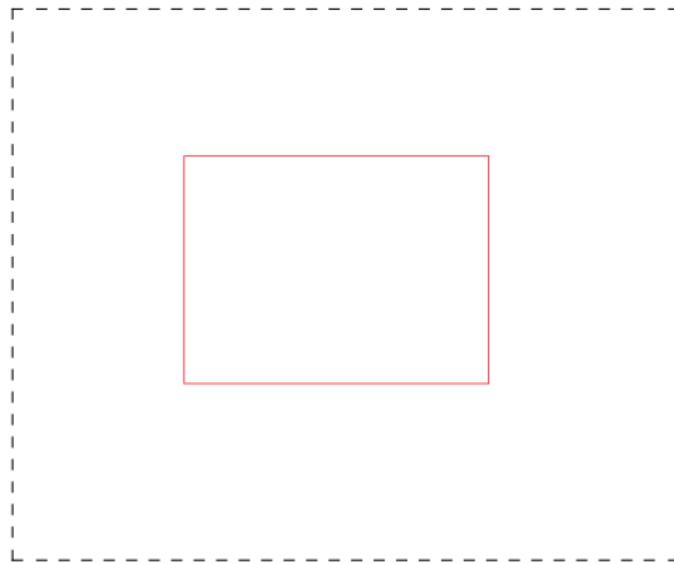


Figure: Dashed line denotes the border of the image (it's not part of the image)

Even worse for synthetic images

- I want to resize my image by factor 2
- And I take every other column and every other row (1st, 3rd, 5th, etc)
- Where is the rectangle!



Figure: Dashed line denotes the border of the image (it's not part of the image)

Even worse for synthetic images

- What's in the image?
- Now I want to resize my image by half in the width direction
- And I take every other column (1st, 3rd, 5th, etc)



Even worse for synthetic images

- What's in the image?
- Now I want to resize my image by half in the width direction
- And I take every other column (1st, 3rd, 5th, etc)

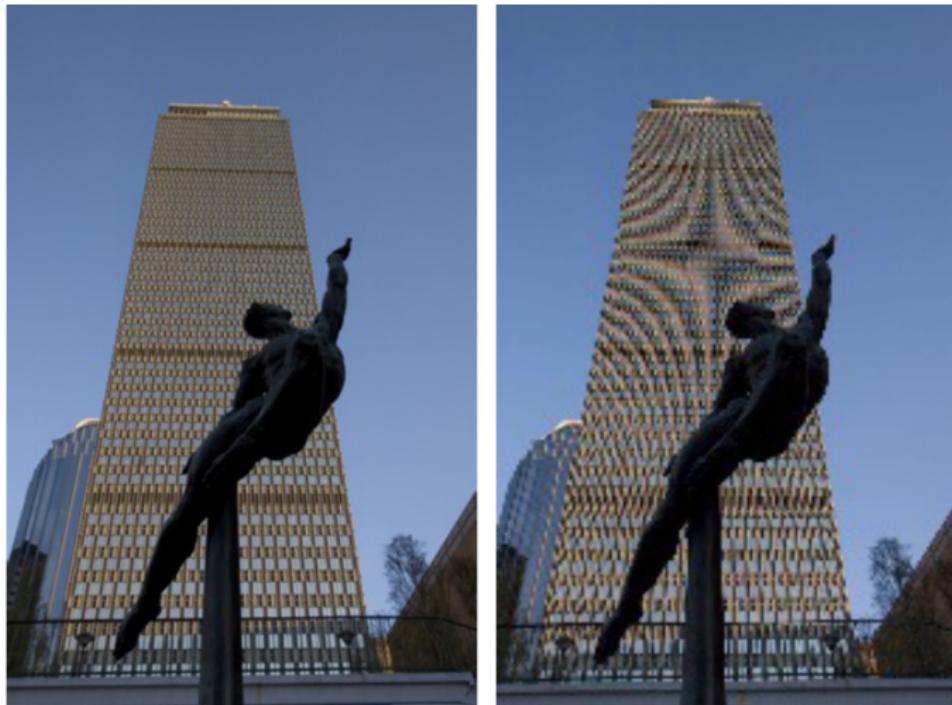


Even worse for synthetic images

- What's in the image?
- Now I want to resize my image by half in the width direction
- And I take every other column (1st, 3rd, 5th, etc)
- Where is the chicken!



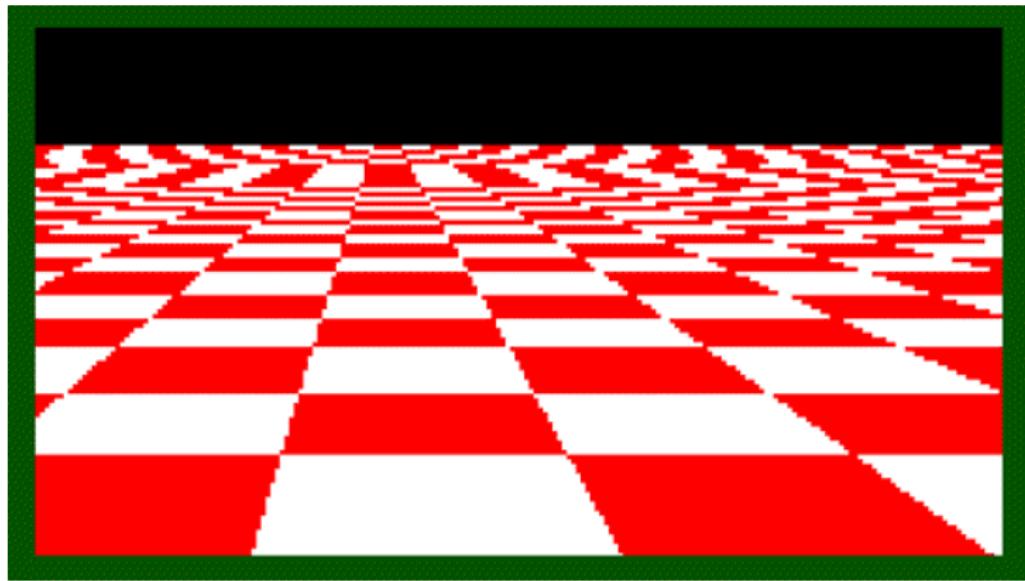
Image Sub-Sampling



[Source: F. Durand]

Even worse for synthetic images

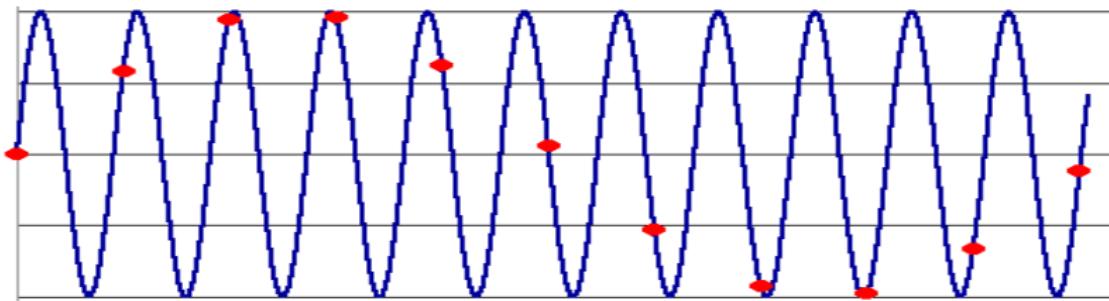
- What's happening?



[Source: L. Zhang]

Aliasing

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image

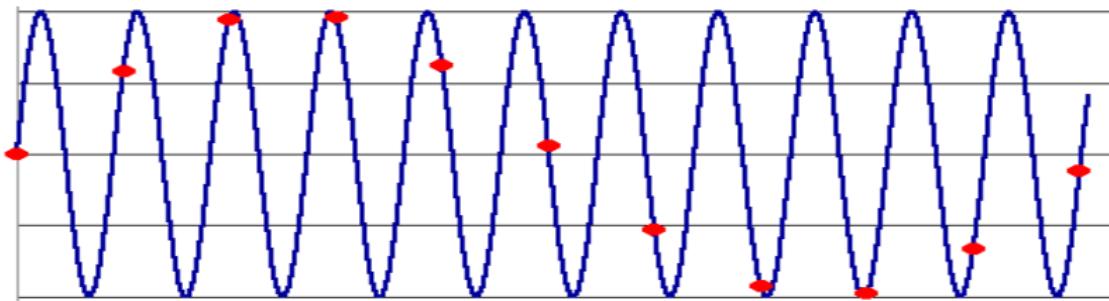


- To do sampling right, need to understand the structure of your signal/image

[Source: R. Urtasun]

Aliasing

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image

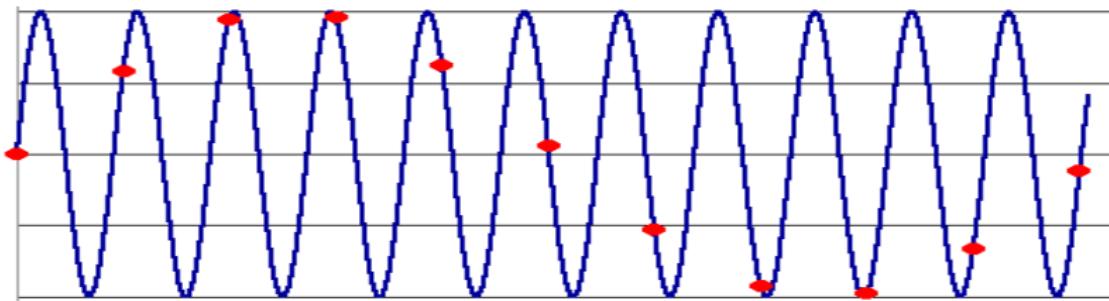


- To do sampling right, need to understand the structure of your signal/image
- The minimum sampling rate is called the **Nyquist rate**

[Source: R. Urtasun]

Aliasing

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image



- To do sampling right, need to understand the structure of your signal/image
- The minimum sampling rate is called the **Nyquist rate**

[Source: R. Urtasun]

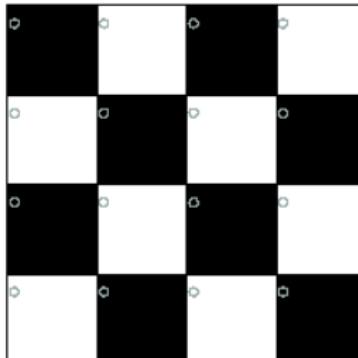
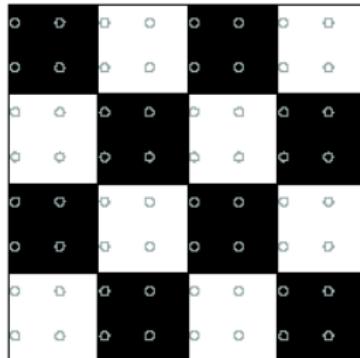
Mr. Nyquist

- Harry Nyquist says that one should look at the frequencies of the signal.
- One should find the highest frequency (via Fourier Transform)
- To sample properly you need to sample with at least twice that frequency
- For those interested: http://en.wikipedia.org/wiki/Nyquist%20Shannon_sampling_theorem

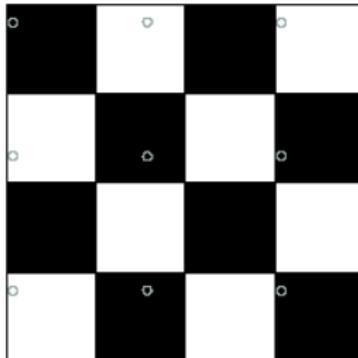
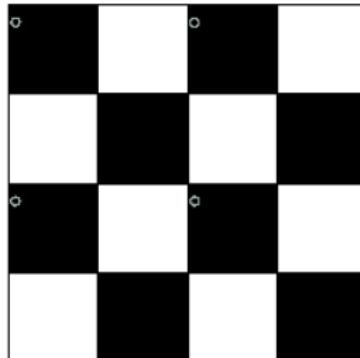
- He looks like a smart guy, we'll just believe him



2D example



Good sampling



Bad sampling

[Source: N. Snavely]

Going back to Downsampling ...

- When downsampling by a factor of two, the original image has frequencies that are too high
- High frequencies are caused by sharp edges
- How can we fix this?

[Adopted from: R. Urtasun]

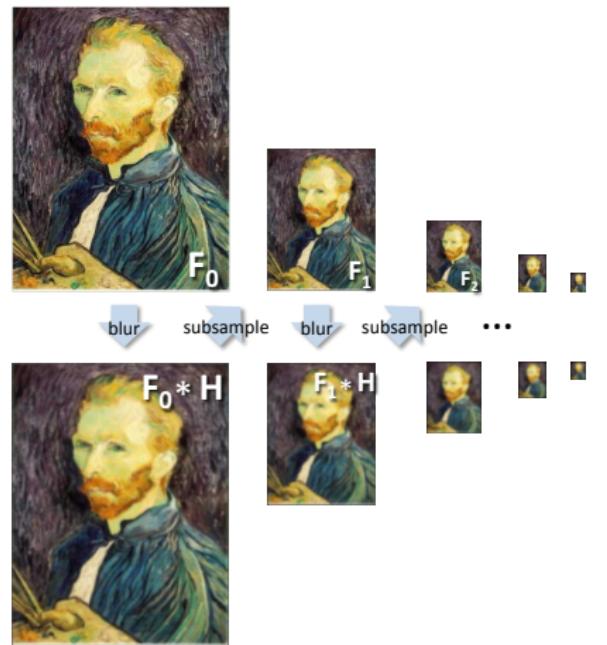
Going back to Downsampling ...

- When downsampling by a factor of two, the original image has frequencies that are too high
- High frequencies are caused by sharp edges
- How can we fix this?

[Adopted from: R. Urtasun]

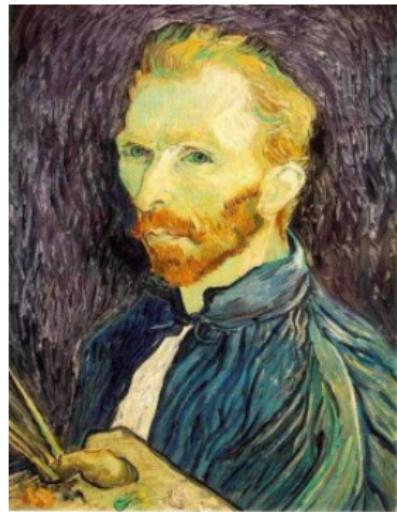
Gaussian pre-filtering

- Solution: Filter out the higher frequency data. Blur the image via Gaussian, then subsample. Very simple!



[Source: N. Snavely]

Subsampling with Gaussian pre-filtering



Gaussian 1/2



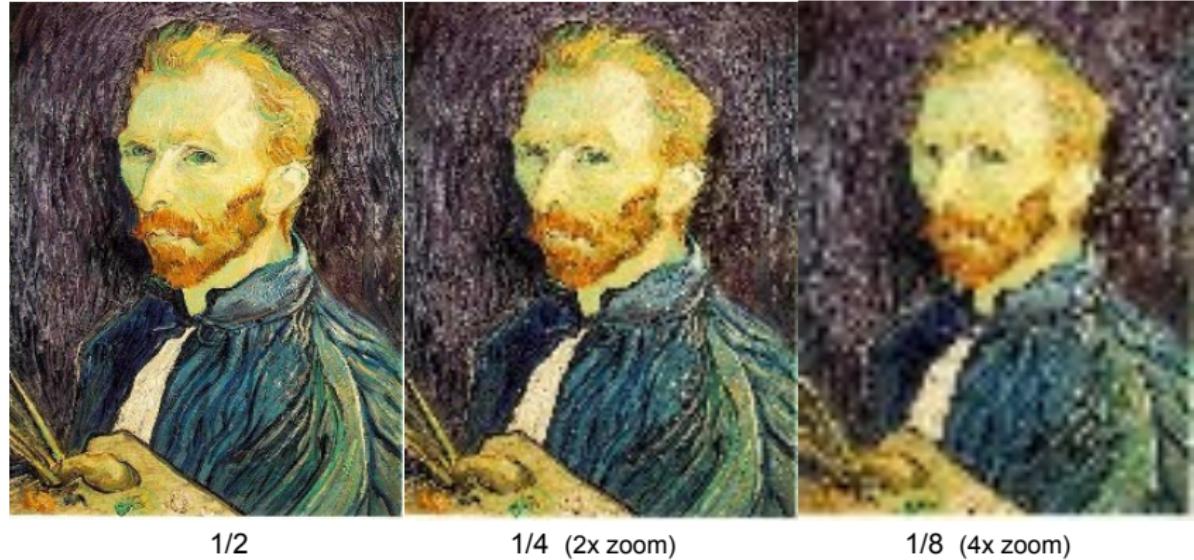
G 1/4



G 1/8

[Source: S. Seitz]

Compare to our result without



[Source: S. Seitz]

Where is the Rectangle?

- My image

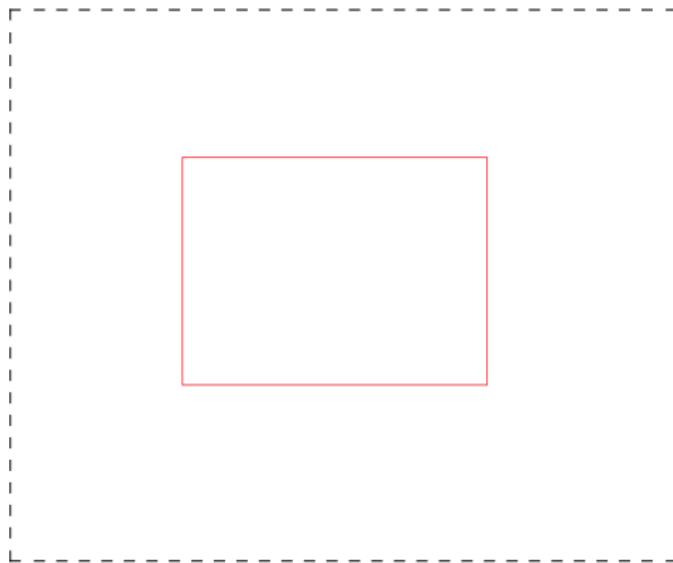


Figure: Dashed line denotes the border of the image (it's not part of the image)

Where is the Rectangle?

- My image
- Let's blur

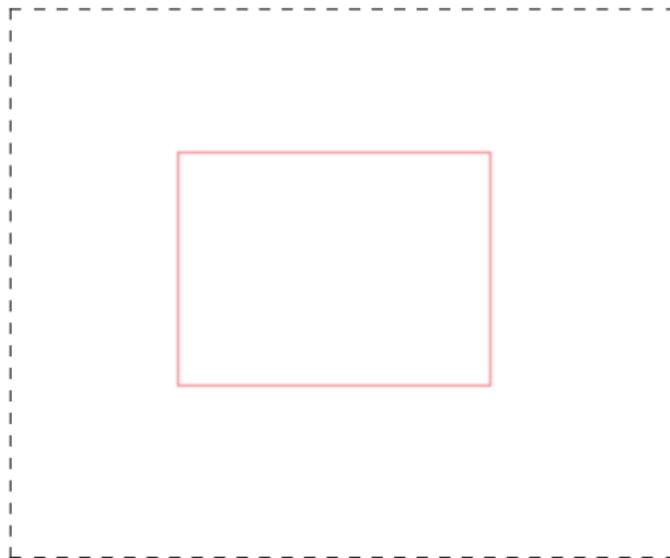


Figure: Dashed line denotes the border of the image (it's not part of the image)

Where is the Rectangle?

- My image
- Let's blur
- And now take every other row and column

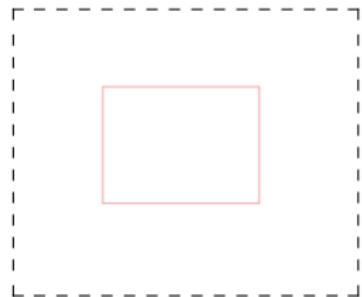


Figure: Dashed line denotes the border of the image (it's not part of the image)

Where is the Chicken?

- My image



Where is the Chicken?

- My image
- Let's blur



Where is the Chicken?

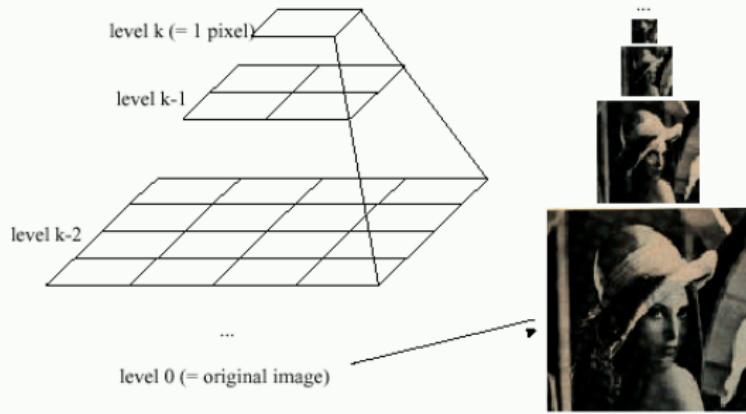
- My image
- Let's blur
- And now take every other column



Gaussian Pyramids [Burt and Adelson, 1983]

- A sequence of images created with Gaussian blurring and downsampling is called a **Gaussian Pyramid**
- In computer graphics, a *mip map* [Williams, 1983]

Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N = 2^k$)



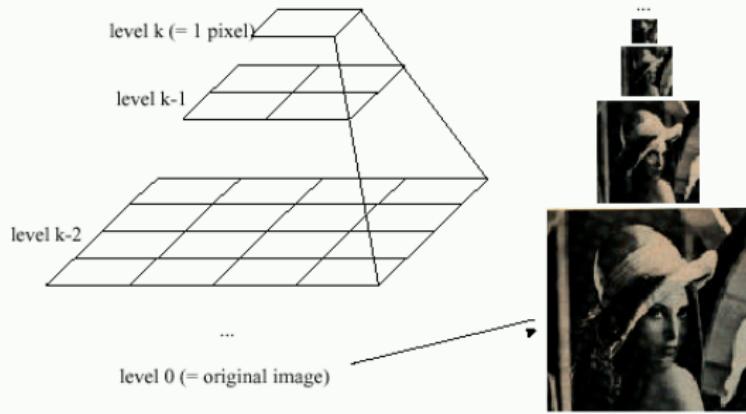
- How much space does a Gaussian pyramid take compared to original image?

[Source: S. Seitz]

Gaussian Pyramids [Burt and Adelson, 1983]

- A sequence of images created with Gaussian blurring and downsampling is called a **Gaussian Pyramid**
- In computer graphics, a *mip map* [Williams, 1983]

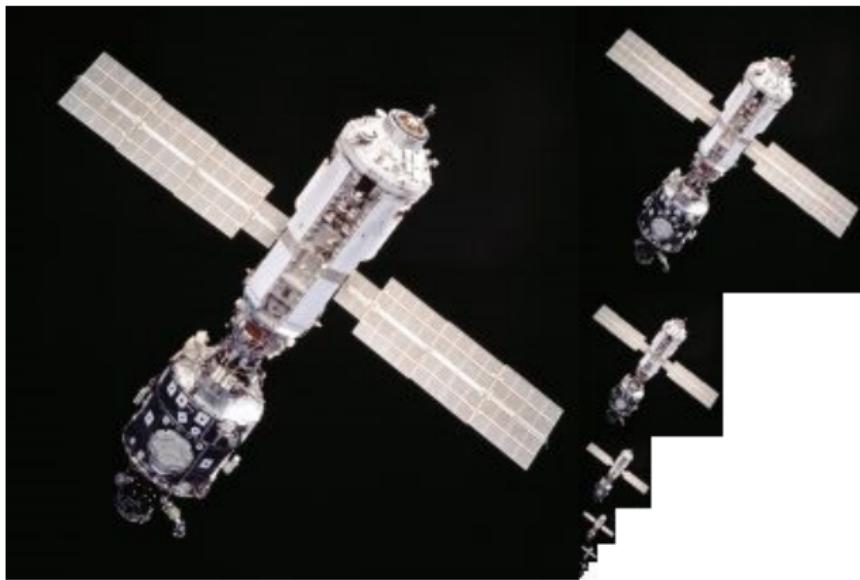
Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N = 2^k$)



- How much space does a Gaussian pyramid take compared to original image?

[Source: S. Seitz]

Example of Gaussian Pyramid



[Source: N. Snavely]

Image Up-Sampling

- This image is too small, how can we make it 10 times as big?



[Source: N. Snavely, R. Urtasun]

Image Up-Sampling

- This image is too small, how can we make it 10 times as big?

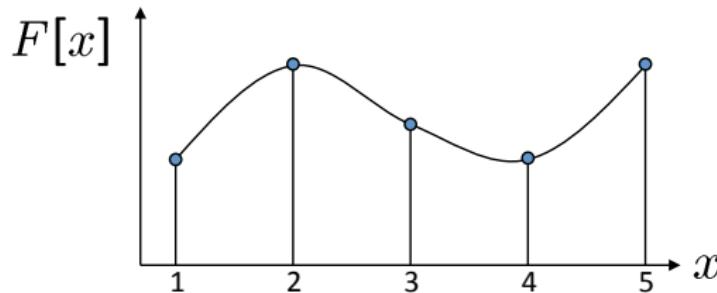


- Simplest approach: repeat each row and column 10 times



[Source: N. Snavely, R. Urtasun]

Interpolation



$d = 1$ in this example

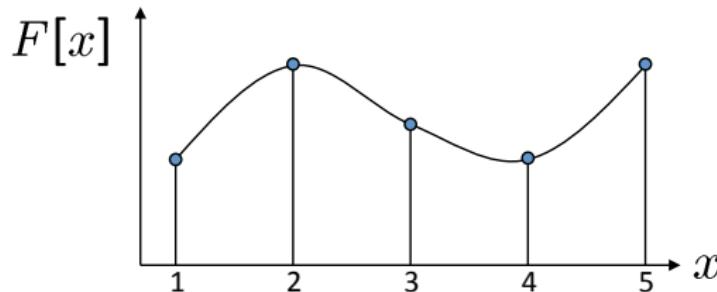
Recall how a digital image is formed

$$F[x, y] = \text{quantize}\left\{f\left(\frac{x}{d}, \frac{y}{d}\right)\right\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

[Source: N. Snavely, S. Seitz]

Interpolation



$d = 1$ in this example

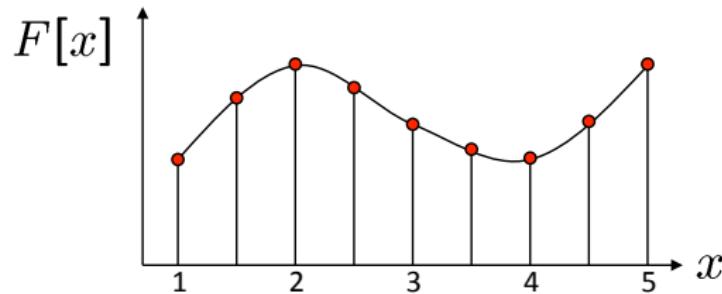
Recall how a digital image is formed

$$F[x, y] = \text{quantize}\left\{f\left(\frac{x}{d}, \frac{y}{d}\right)\right\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

[Source: N. Snavely, S. Seitz]

Interpolation



$d = 1$ in this example

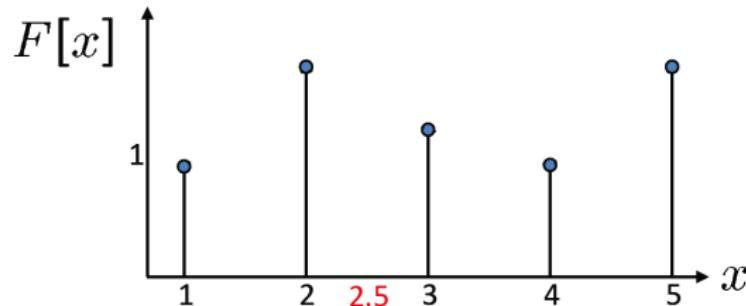
Recall how a digital image is formed

$$F[x, y] = \text{quantize}\left\{f\left(\frac{x}{d}, \frac{y}{d}\right)\right\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

[Source: N. Snavely, S. Seitz]

Interpolation

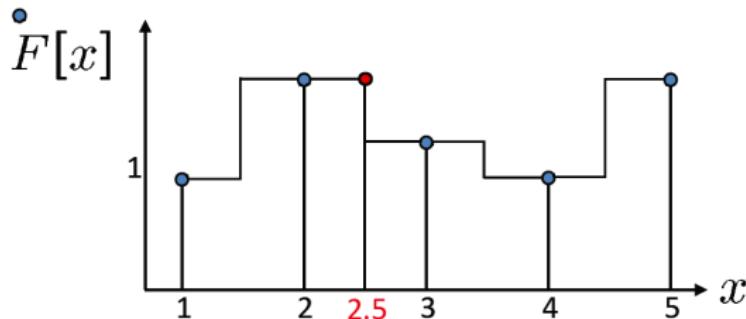


$d = 1$ in this example

What if we don't know f ?

[Source: N. Snavely, S. Seitz]

Interpolation



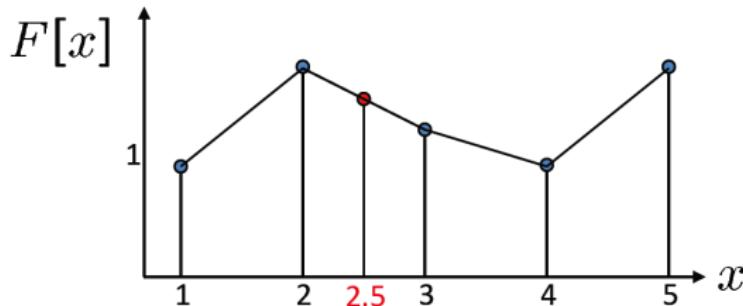
$d = 1$ in this example

What if we don't know f ?

- Guess an approximation: for example nearest-neighbor

[Source: N. Snavely, S. Seitz]

Interpolation



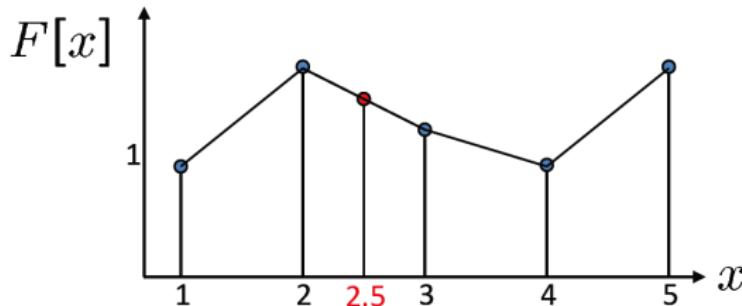
$d = 1$ in this example

What if we don't know f ?

- Guess an approximation: for example nearest-neighbor
- Guess an approximation: for example linear

[Source: N. Snavely, S. Seitz]

Interpolation



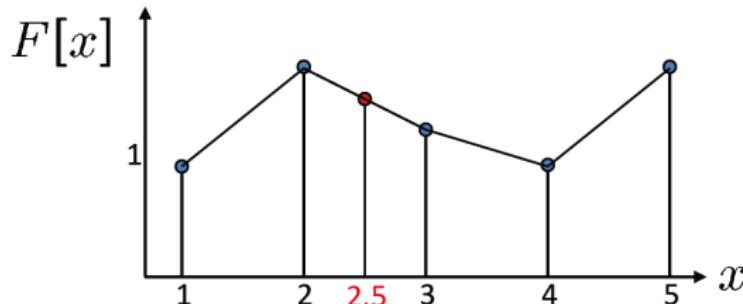
$d = 1$ in this example

What if we don't know f ?

- Guess an approximation: for example nearest-neighbor
- Guess an approximation: for example linear
- More complex approximations: cubic, B-splines

[Source: N. Snavely, S. Seitz]

Interpolation



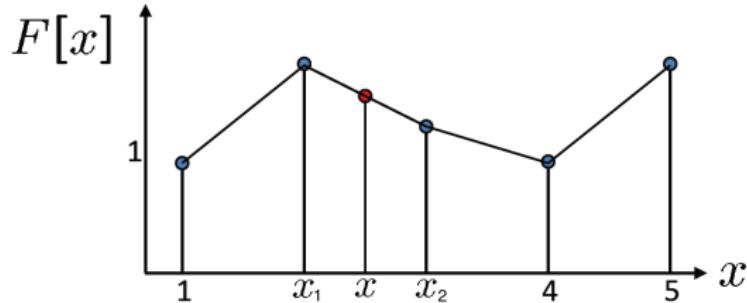
$d = 1$ in this example

What if we don't know f ?

- Guess an approximation: for example nearest-neighbor
- Guess an approximation: for example linear
- More complex approximations: cubic, B-splines
- But more isn't always better!

[Source: N. Snavely, S. Seitz]

Linear Interpolation



$d = 1$ in this example

- Linear interpolation from our discretized F :

$$G(x) = \frac{x_2 - x}{x_2 - x_1} F(x_1) + \frac{x - x_1}{x_2 - x_1} F(x_2)$$

Interpolation: 1D Example



- Let's make this signal triple length

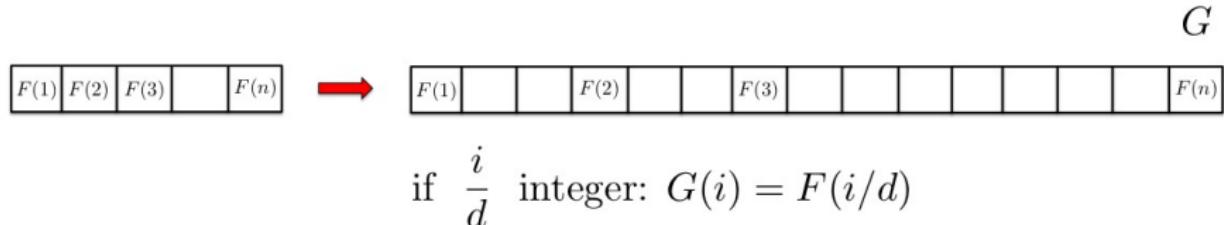
Interpolation: 1D Example



Make a vector G with d times the size of F

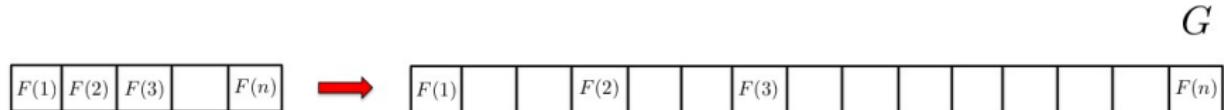
- Let's make this signal triple length ($d = 3$)

Interpolation: 1D Example



- Let's make this signal triple length ($d = 3$)
- If i/d is an integer, just copy from the signal

Interpolation: 1D Example



if $\frac{i}{d}$ integer: $G(i) = F(i/d)$

otherwise: $G(i) = \frac{x_2 - x}{x_2 - x_1}F(x_1) + \frac{x - x_1}{x_2 - x_1}F(x_2)$

$$\begin{aligned} & x = i/d \\ \text{where } & x_1 = \lfloor i/d \rfloor \\ & x_2 = \lceil i/d \rceil \end{aligned}$$

- Let's make this signal triple length ($d = 3$)
- If i/d is an integer, just copy from the signal
- Otherwise use the interpolation formula

Interpolation via Convolution: 1D Example



- Let's make this signal triple length

Interpolation via Convolution: 1D Example

$F(1)$	$F(2)$	$F(3)$		$F(n)$
--------	--------	--------	--	--------



$F(1)$	0	0	$F(2)$	0	0	$F(3)$	0	0					$F(n)$
--------	---	---	--------	---	---	--------	---	---	--	--	--	--	--------

if $\frac{i}{d}$ integer: $G'(i) = F(i/d)$

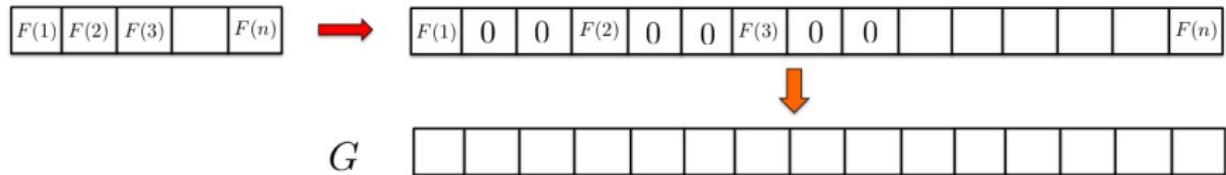
otherwise: 0

- Let's make this signal triple length ($d = 3$)

Interpolation via Convolution: 1D Example

$h = ?$

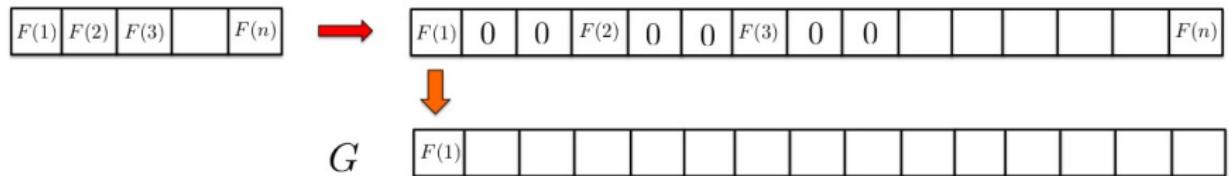
$* G'$



- Let's make this signal triple length ($d = 3$)
- What should be my "reconstruction" filter h (such that $G = h * G'$)?

Interpolation via Convolution: 1D Example

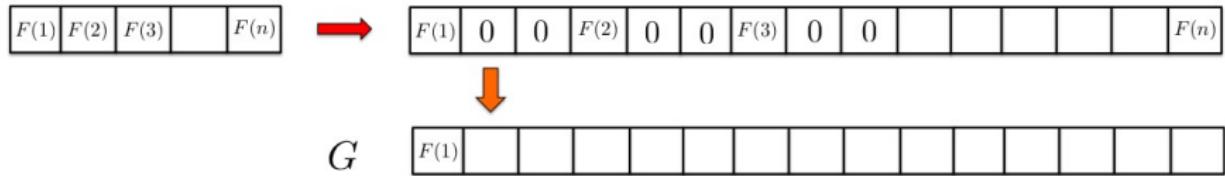
$$h = \left[0, \frac{1}{3}, \frac{2}{3}, 1, \frac{2}{3}, \frac{1}{3}, 0\right] * G'$$



- Let's make this signal triple length ($d = 3$)
- What should be my “reconstruction” filter h (such that $G = h * G'$)?
- $h = [0, \frac{1}{d}, \dots, \frac{d-1}{d}, 1, \frac{d-1}{d}, \dots, \frac{1}{d}, 0]$, where d my upsampling factor

Interpolation via Convolution: 1D Example

$$h = \left[0, \frac{1}{3}, \frac{2}{3}, 1, \frac{2}{3}, \frac{1}{3}, 0\right] * G'$$

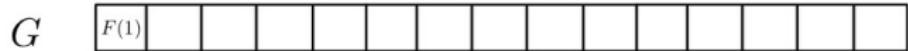
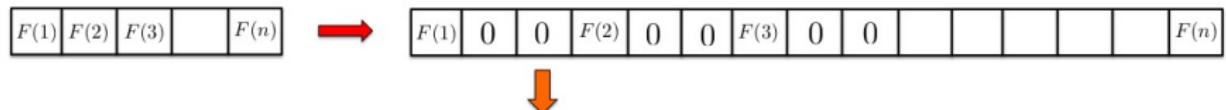


$$\frac{2}{3}F(1) + \frac{1}{3}F(2)$$

- Let's make this signal triple length ($d = 3$)
- What should be my “reconstruction” filter h (such that $G = h * G'$)?
- $h = [0, \frac{1}{d}, \dots, \frac{d-1}{d}, 1, \frac{d-1}{d}, \dots, \frac{1}{d}, 0]$, where d my upsampling factor

Interpolation via Convolution: 1D Example

$$h = \left[0, \frac{1}{3}, \frac{2}{3}, 1, \frac{2}{3}, \frac{1}{3}, 0\right] * G'$$

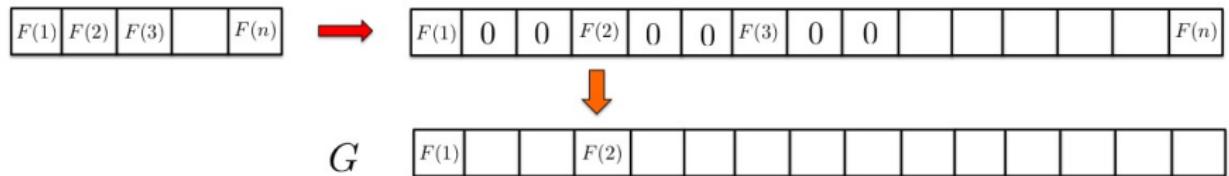


$$\frac{1}{3}F(1) + \frac{2}{3}F(2)$$

- Let's make this signal triple length ($d = 3$)
- What should be my “reconstruction” filter h (such that $G = h * G'$)?
- $h = [0, \frac{1}{d}, \dots, \frac{d-1}{d}, 1, \frac{d-1}{d}, \dots, \frac{1}{d}, 0]$, where d my upsampling factor

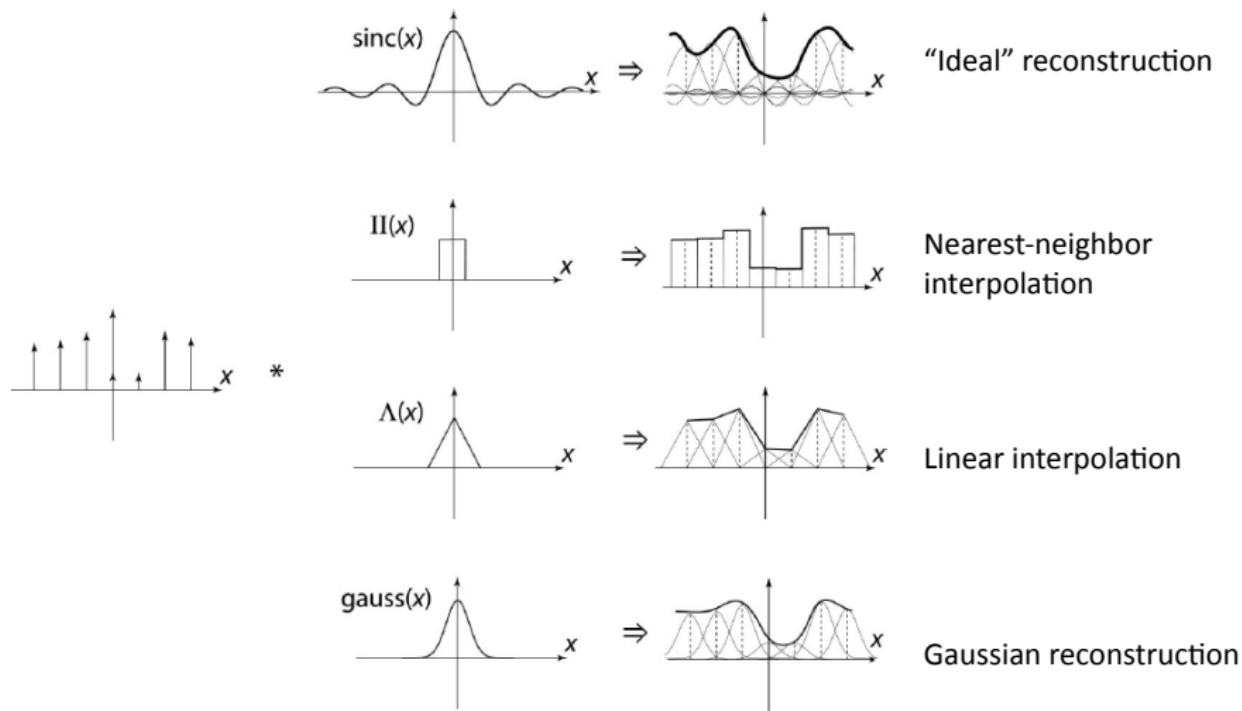
Interpolation via Convolution: 1D Example

$$h = \left[0, \frac{1}{3}, \frac{2}{3}, 1, \frac{2}{3}, \frac{1}{3}, 0\right] * G'$$



- Let's make this signal triple length ($d = 3$)
- What should be my “reconstruction” filter h (such that $G = h * G'$)?
- $h = [0, \frac{1}{d}, \dots, \frac{d-1}{d}, 1, \frac{d-1}{d}, \dots, \frac{1}{d}, 0]$, where d my upsampling factor

Interpolation via Convolution (1D)



Source: B. Curless

Image Interpolation (2D)

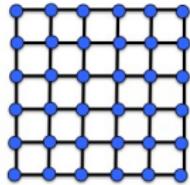
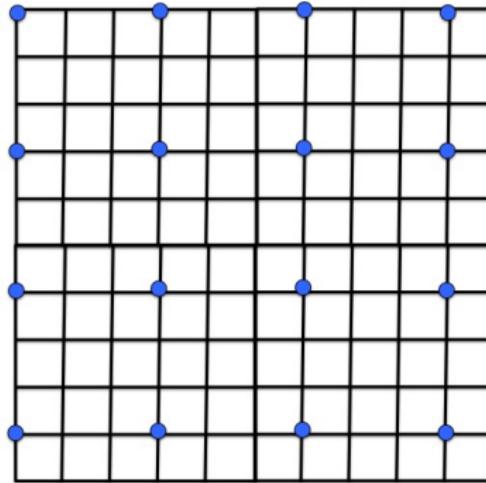
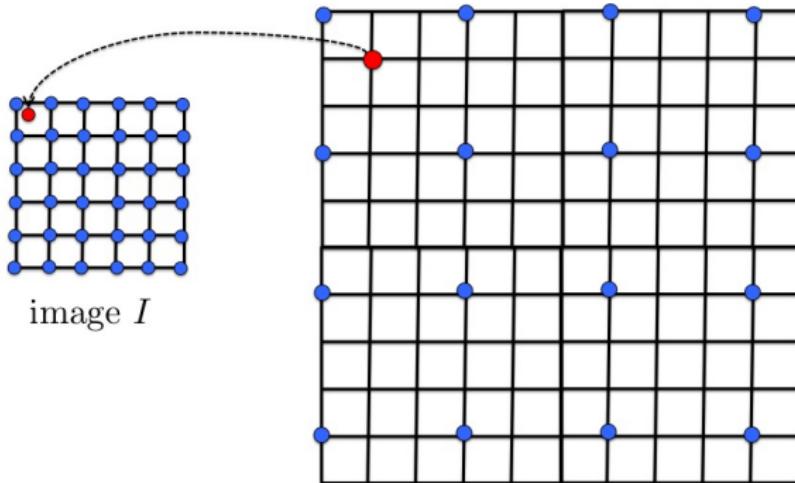


image I



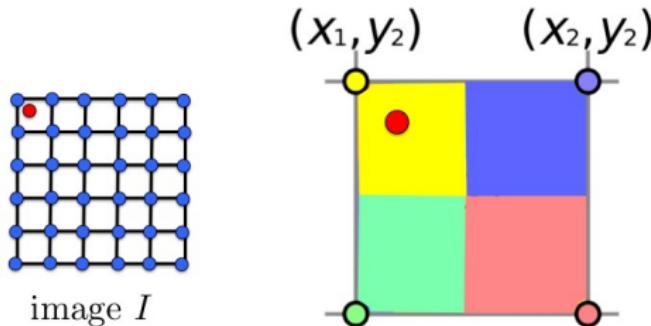
- Let's make this image triple size
- Copy image in every third pixel. What about the remaining pixels in G ?

Image Interpolation (2D)



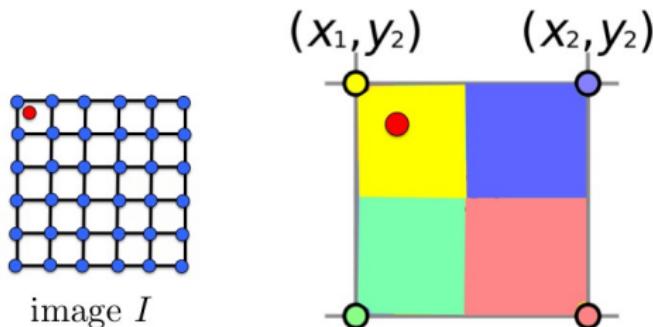
- Let's make this image triple size
- Copy image in every third pixel. What about the remaining pixels in G ?
- How shall we compute this value?

Image Interpolation (2D)



- Let's make this image triple size
- Copy image in every third pixel. What about the remaining pixels in G ?
- One possible way: nearest neighbor interpolation

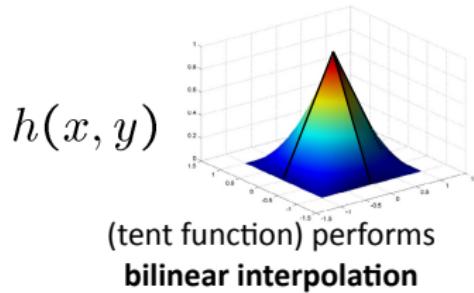
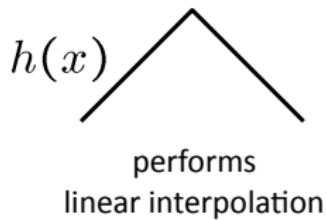
Image Interpolation (2D)



- Let's make this image triple size
- Copy image in every third pixel. What about the remaining pixels in G ?
- Better: bilinear interpolation; linear interpolation in x , then y , resulting in a quadratic interpolation.
- Check out details:
http://en.wikipedia.org/wiki/Bilinear_interpolation

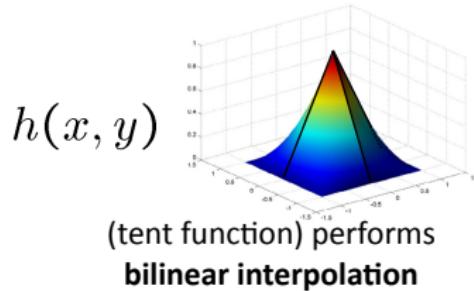
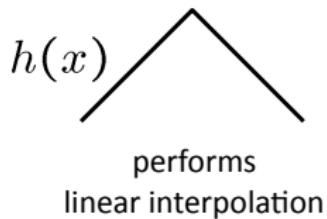
Reconstruction Filters

- What does the 2D version of this hat function look like?



Reconstruction Filters

- What does the 2D version of this hat function look like?



- And filter for nearest neighbor interpolation?

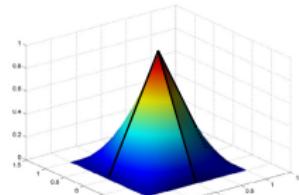
Reconstruction Filters

- What does the 2D version of this hat function look like?



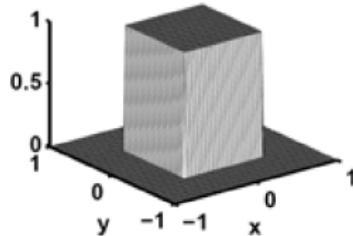
performs
linear interpolation

$$h(x, y)$$



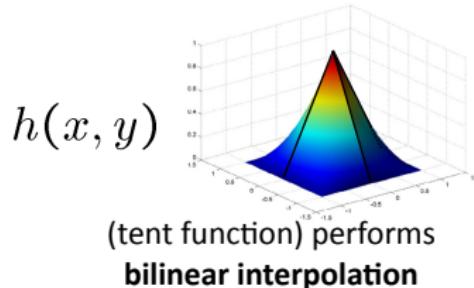
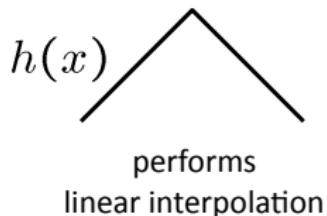
(tent function) performs
bilinear interpolation

- And filter for nearest neighbor interpolation?



Reconstruction Filters

- What does the 2D version of this hat function look like?



- Better filters give better resampled images: Bicubic is a common choice

Image Interpolation via Convolution (2D)

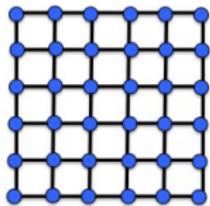
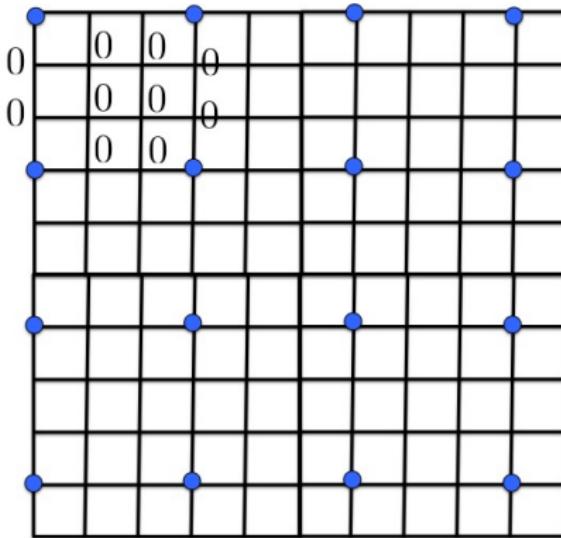
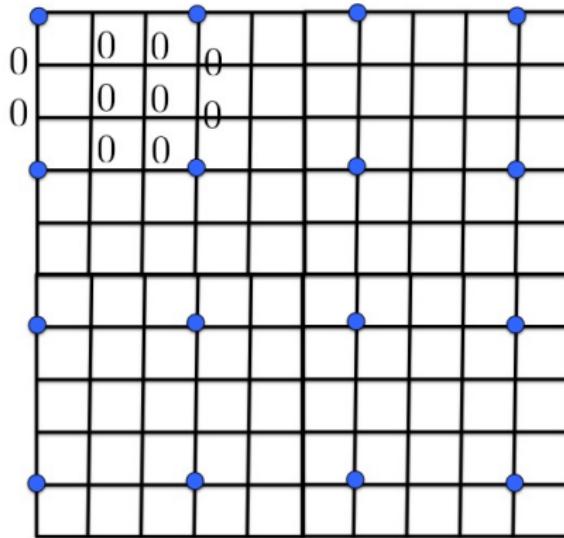
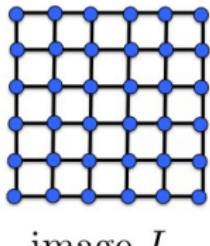


image I

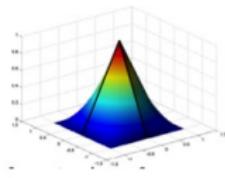


- Let's make this image triple size: copy image values in every third pixel, place zeros everywhere else

Image Interpolation via Convolution (2D)



$$\ast h(x, y)$$



- Let's make this image triple size: copy image values in every third pixel, place zeros everywhere else
- Convolution with a reconstruction filter (e.g., bilinear) and you get the interpolated image

Image Interpolation

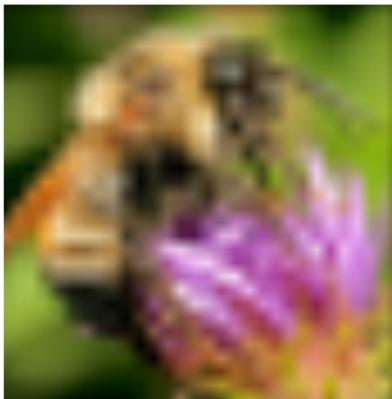
Original image



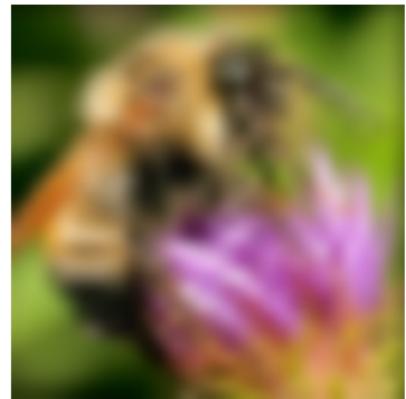
Interpolation results



Nearest-neighbor interpolation



Bilinear interpolation



Bicubic interpolation

[Source: N. Snavely]

Summary – Stuff You Should Know

- To down-scale an image: blur it with a small Gaussian (e.g., $\sigma = 1.4$) and downsample
- To up-scale an image: interpolation (nearest neighbor, bilinear, bicubic, etc)
- Gaussian pyramid: Blur with Gaussian filter, downsample result by factor 2, blur it with the Gaussian, downsample by 2...

Matlab functions:

- FSPECIAL: creates a Gaussian filter with specified σ
- IMFILTER: convolve image with the filter
- I(1:2:END, 1:2:END): takes every second row and column
- IMRESIZE(IMAGE, SCALE, METHOD): Matlab's function for resizing the image, where METHOD=“nearest”, “bilinear”, “bicubic” (works for downsampling and upsampling)